

kartik Thakur

Assignment 19

Predict whether the person is having diabetes or not

1. Logistic Regression

2. Decision Tree

3. Random Forest

4. Adaptive Boosting

5. Gradient Boosting

Compare the results and tell which algorithm is good for predicting the person is diabetic or not

1. Logistic Regression

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [2]: db=pd.read_csv("diabetes - diabetes.csv")
```

```
In [3]: db.head()
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6.0	148.0	72.0	35.0	0.0	33.6	0.627
1	1.0	85.0	66.0	29.0	0.0	26.6	0.351
2	8.0	183.0	64.0	0.0	0.0	23.3	0.672
3	1.0	89.0	66.0	23.0	94.0	28.1	0.167
4	0.0	137.0	40.0	35.0	168.0	43.1	2.288

In [6]: `db.describe()`

Out[6]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
count	767.000000	759.000000	741.000000	747.000000	732.000000	746.000000	
mean	3.844850	121.001318	69.026991	20.416332	78.285519	32.013673	
std	3.371772	32.071511	19.378222	15.976672	113.071200	7.896529	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.325000	
50%	3.000000	117.000000	72.000000	23.000000	22.500000	32.050000	
75%	6.000000	141.000000	80.000000	32.000000	126.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In [7]: `db.isna().sum()`

Out[7]:

Pregnancies	1
Glucose	9
BloodPressure	27
SkinThickness	21
Insulin	36
BMI	22
DiabetesPedigreeFunction	16
Age	6
Outcome	0
dtype: int64	

In [8]:

```

db['Pregnancies'].fillna(db['Pregnancies'].mean(),inplace=True)
db['Glucose'].fillna(db['Glucose'].mean(),inplace=True)
db['BloodPressure'].fillna(db['BloodPressure'].mean(),inplace=True)
db['SkinThickness'].fillna(db['SkinThickness'].mean(),inplace=True)
db['Insulin'].fillna(db['Insulin'].mean(),inplace=True)
db['BMI'].fillna(db['BMI'].mean(),inplace=True)
db['DiabetesPedigreeFunction'].fillna(db['DiabetesPedigreeFunction'].mean(),inplace=True)
db['Age'].fillna(db['Age'].mean(),inplace=True)

```

In [9]: `db.isna().sum()`

Out[9]:

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	0
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	0
Outcome	0
dtype: int64	

```
In [10]: db.columns
```

```
Out[10]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

```
In [11]: x=db[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age']]
         y=db['Outcome']
```

```
In [12]: from sklearn.model_selection import train_test_split
```

```
In [13]: x_train,x_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=
```

```
In [14]: from sklearn.linear_model import LogisticRegression
```

```
In [15]: model= LogisticRegression()
```

```
In [53]: model.fit(x_train,y_train)
```

C:\Users\keerti chouhan\anaconda3\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[53]: LogisticRegression()
```


```
In [17]: predictions=model.predict(x_test)
```

```
In [18]: predictions
```

```
Out[18]: array([1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0,
                0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1,
                0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1,
                0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1,
                1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 0,
                0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0],
              dtype=int64)
```

```
In [19]: plt.scatter(y_test,predictions)
```

```
Out[19]: <matplotlib.collections.PathCollection at 0x23351dd0850>
```



```
In [20]: from sklearn.metrics import confusion_matrix
         from sklearn.metrics import classification_report
```

```
In [21]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.81	0.87	0.84	100
1	0.72	0.61	0.66	54
accuracy			0.78	154
macro avg	0.76	0.74	0.75	154
weighted avg	0.77	0.78	0.77	154

```
In [22]: print(confusion_matrix(y_test,predictions))
```

```
[[87 13]
 [21 33]]
```

Decision Tree

```
In [23]: from sklearn.tree import DecisionTreeClassifier
```

```
In [24]: dtree=DecisionTreeClassifier()
```

```
In [25]: dtree.fit(x_train,y_train)
```

```
Out[25]: DecisionTreeClassifier()
```


```
In [26]: pridiction1=dtree.predict(x_test)
```

```
In [27]: pridiction1
```

```
Out[27]: array([1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0,
        0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1,
        0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0,
        0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
        1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0,
        1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0,
        0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1],
        dtype=int64)
```

```
In [28]: plt.scatter(y_test,predictions)
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x23351fb6b50>
```



```
In [29]: print(classification_report(y_test,pridiction1))
```

	precision	recall	f1-score	support
0	0.80	0.72	0.76	100
1	0.56	0.67	0.61	54
accuracy			0.70	154
macro avg	0.68	0.69	0.68	154
weighted avg	0.72	0.70	0.71	154

```
In [30]: print(confusion_matrix(y_test,pridiction1))
```

```
[[72 28]
 [18 36]]
```

.Random Forest

```
In [31]: from sklearn.ensemble import RandomForestClassifier
```

```
In [32]: rfc=RandomForestClassifier(n_estimators=600)
```

```
In [33]: rfc.fit(x_train,y_train)
```

```
Out[33]: RandomForestClassifier(n_estimators=600)
```

```
In [34]: pridiction2=rfc.predict(x_test)
```

```
In [35]: pridiction2
```

```
Out[35]: array([1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0,
                0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 1, 1,
                0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 0,
                1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
                1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0],
                dtype=int64)
```

```
In [36]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [37]: print(classification_report(y_test,pridiction2))
```

	precision	recall	f1-score	support
0	0.82	0.88	0.85	100
1	0.74	0.65	0.69	54
accuracy			0.80	154
macro avg	0.78	0.76	0.77	154
weighted avg	0.80	0.80	0.80	154

```
In [38]: print(confusion_matrix(y_test,pridiction2))
```

```
[[88 12]
 [19 35]]
```

4.Adaptive Boosting

```
In [39]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [40]: Ada= AdaBoostClassifier(n_estimators=150)
```

```
In [41]: Ada.fit(x_train,y_train)
```

```
Out[41]: AdaBoostClassifier(n_estimators=150)
```

```
In [42]: pridiction3=Ada.predict(x_test)
```

```
In [43]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [44]: print(classification_report(y_test,pridiction3))
```

	precision	recall	f1-score	support
0	0.75	0.76	0.75	100
1	0.54	0.52	0.53	54
accuracy			0.68	154
macro avg	0.64	0.64	0.64	154
weighted avg	0.67	0.68	0.67	154

```
In [45]: print(confusion_matrix(y_test,pridiction3))
```

```
[[76 24]
 [26 28]]
```

5Gradient Boosting

```
In [46]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [47]: GD= AdaBoostClassifier(n_estimators=200)
```

```
In [48]: GD.fit(x_train,y_train)
```

```
Out[48]: AdaBoostClassifier(n_estimators=200)
```

```
In [49]: pridiction4=GD.predict(x_test)
```



```
In [50]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [51]: print(classification_report(y_test, pridiction4))
```

	precision	recall	f1-score	support
0	0.75	0.78	0.76	100
1	0.56	0.52	0.54	54
accuracy			0.69	154
macro avg	0.66	0.65	0.65	154
weighted avg	0.68	0.69	0.69	154

```
In [ ]:
```